

Constructing a Zero-Trust Kubernetes Cluster

Bachelor Thesis

Submitted in partial fulfillment of the requirements for the degree of

Bachelor of Science in Engineering

to the University of Applied Sciences FH Campus Wien

Bachelor Degree Program: Computer Science and Digital Communications

Author:

Guntram Björn Klaus

Student identification number:

c2110475170

Supervisor:

BSc. MSc. Bernhard Taufner

Date:

dd.mm.yyyy

Declaration of authorship:

I declare that this thesis is my own work and that I did not use any aids other than those indicated or any other unauthorized help (e.g., ChatGPT or similar artificial intelligence-based programs). I certify that this work does not contain any personal data, and that I have clarified any copyright, license or image-law issues pertaining to the electronic publication of this thesis. Otherwise, I will indemnify and hold harmless the FH Campus Wien from any claims for compensation by third parties. I certify that I have not submitted this thesis (to an assessor for review) in Austria or abroad in any form as an examination paper. I further certify that the (printed and electronic) copies I have submitted are identical.

Date:

Signature:

Abstract

(E.g. “This thesis investigates...”)

Kurzfassung

(Z.B. "Diese Arbeit untersucht...")

List of Abbreviations

AKS	Azure Kubernetes Service
EKS	Elastic Kubernetes Service
FOSS	Free and Open Source Software
GKE	Google Kubernetes Engine
IT	Information Technology
IoT	Internet of Things
MFA	Multifactor Authentication
mTLS	Mutual Transport Layer Security
NIST	National Institute of Standards and Technology
VPN	Virtual Private Network
ZT	Zero Trust

Key Terms

Kubernetes

Cloud

Zero Trust

Least Privilege

Access Control

Contents

1	Introduction	1
1.1	Kubernetes and Zero Trust	1
1.2	Related Work	1
1.3	Objectives and Methodology	3
2	Concepts	4
2.1	Zero Trust	4
2.1.1	Definition and principles	4
2.1.2	The network	4
2.1.3	Logical components	5
2.2	Kubernetes	6
2.2.1	Components	6
2.2.2	Core Cluster Objects	7
2.2.3	Kubernetes Networking	8
3	Creating the Zero Trust Cluster	10
3.1	Virtual Machines and Network	10
3.2	Bootstrapping the cluster	11
3.3	Teleport	14
3.3.1	Why use Teleport?	14
3.3.2	Installation and setup	16
3.4	Istio	18
3.4.1	Service Meshes	18
3.4.2	mTLS	18
3.4.3	Why use Istio?	18
3.4.4	Installation and setup	20
3.5	Keycloak	20
3.5.1	Identity and Access Management	20
3.5.2	RBAC and ABAC	20
3.5.3	OIDC and OAuth2	21
3.5.4	Why use Keycloak?	21
3.5.5	Installation and setup	22
3.6	Vault Secret Management	22
3.6.1	Why use Vault?	22
3.7	Trivy	22
3.7.1	Why Trivy?	22
4	Discussion and Future Work	23
5	Conclusion	24
	Bibliography	25
	List of Figures	26

List of Tables	27
Appendix	28

1 Introduction

1.1 Kubernetes and Zero Trust

Over the last years, there have been two significant shifts in enterprise IT systems.

One of these shifts addresses the way companies deploy, scale and maintain the lifecycle of their software services. Applications used to be primarily constructed as one large software unit that bundled all features, business logic, user interfaces and data access components. Increasing necessity for scalability, flexibility and maintainability made organizations transition from such monolithic architectures to microservices, which are smaller, separated but loosely coupled software units implementing one component of the larger system at hand. The release of the Docker container platform in 2013 had a significant impact on making transitions from monoliths to microservices feasible. Kubernetes has since emerged as the go-to choice for managing containerized applications at scale, particularly in cloud-native environments. Its ability to automate tasks, scale applications, and support a wide range of use cases makes it a powerful tool for both large enterprises and small teams.

The second shift pertains to how companies secure their computing infrastructure and the resource hosted on it. While there used to be single, easily identifiable network perimeters in the past, for example a single local area network at a company site, modern infrastructures may consist of multiple internal networks, remote offices, mobile workers, different types of virtualization and cloud services. This circumstance has rendered traditional, static, perimeter-based security no longer appropriate, because transgressing this perimeter once means further, unhindered access into a given system. Under a newer security model labeled "Zero Trust", the aim is to restrict such unhindered movement as much as possible by adhering to certain principles and guidelines, the central one being to never grant implicit trust to any actor on a network - hence the term "Zero" Trust. The principles of ZT have existed way before the term "Zero Trust" was coined. A fundamental piece of literature, which tries to capture what "Zero Trust" concretely means, is the NIST Special Publication 800-207, titled "Zero Trust Architecture". It is the point of reference for understanding Zero Trust.

As both topics of container orchestration with Kubernetes and improved security paradigms are evermore gaining traction in enterprise systems, it is of paramount importance to explore how exactly this current new paradigm of Zero Trust can be applied to Kubernetes.

1.2 Related Work

Since the emergence and popularization of the concept of "Zero-Trust", a lot of work has been done on the topic, tying it into various domains of IT: Cloud, on-premise infrastructure, IoT, Hardware, Blockchain and much more.

Andrea Manzato, at the University of Padua, implements the Zero Trust model in an enterprise environment using solutions provided by Microsoft Azure. It is investigated how

the capabilities and configuration options of Microsoft Defender and Active Directory can be leveraged to protect enterprise resources. Attack scenarios on these technologies are simulated and automated remediation actions are presented.

Dr. Wesam Almobaideen's master thesis, at Rochester Institute of Technology Dubai Campus, explores the topic of Zero-Trust specifically in the context of MFA. A framework combining principles of ZT and MFA is designed and evaluated in terms of performance, security and user satisfaction.

In the context of IoT, Cem Bicer at the technical university of Vienna, explores and evaluates ZT for edge networks. The implementation of the thesis follows ZT guidelines as proposed by the National Institute of Standards and Technology (NIST) and additionally places a blockchain network on top of the ZT architecture.

Furthermore, zero trust in and of itself has been put under scrutiny. In "Theory and Application of Zero Trust Security: A Brief Survey", Kang et. al investigates the current challenges faced when making use of Zero Trust, as well as progress that has been achieved so far when doing so. Here, it is noted that research and knowledge on the theory and application of Zero Trust has not yet matured, and more extensive work is still required to obtain a deeper understanding and more accurate implementation of the paradigm in academia and industry. In his master's thesis at Utrecht university, Michel Modderkolk proposes a more mature model of Zero Trust, branded "Zero Trust Maturity Model (ZeTuMM)". The work makes use of two scientific methods, "comparison analysis" and "focus area maturity modeling", in order to outline a more fully fledged ZT model.

Independent of ZT, Kubernetes security has been explored in the following works. The paper titled "Designing an intrusion detection system for a Kubernetes cluster" by Hristov, P. (2022) focuses on addressing the security challenges associated with cloud-based architectures, particularly those utilizing Kubernetes. The essay highlights the growing reliance on digital platforms, which has led to an increased demand for automation and high-reliability systems. This reliance has, in turn, increased the use of tools and consequently, the number of security concerns.

"A Systematic evaluation of CVEs and mitigation strategies for a Kubernetes stack" by Fred Nordell (2022) offers a detailed examination of CVE's and possible mitigation strategies. The thesis "Kubernetes Near Real-Time Monitoring and Secure Network Architectures" explores the security challenges and mitigation strategies in Kubernetes. It emphasizes the importance of securing the Kubernetes control plane, implementing secure configurations, and protecting critical components like the etcd data store. Additionally, it highlights the integration of security practices throughout the development, deployment, and operations lifecycle, advocating for a DevSecOps approach that includes security automation, continuous security testing, and security monitoring.

"Testing the Security of a Kubernetes Cluster in a Production Environment" by Giangiulio and Malmberg at KTH Stockholm emphasizes the critical importance of comprehensive security measures for Kubernetes clusters operating in production environments. It highlights the need for robust security practices, including data encryption, proper user and permissions management, and the use of Role-Based Access Control (RBAC) to manage access to the cluster. Additionally, it underscores the significance of maintaining up-to-date Docker images, using minimal Docker images to reduce potential vulnerabilities, and implementing network policies and secrets management to enhance security.

The thesis "A Security Framework for Multi-Cluster Kubernetes Architectures" also aims to address the security challenges in Kubernetes environments, particularly focusing on multi-cluster setups. The study emphasizes the importance of continuous monitoring and manage-

ment to mitigate security risks, highlighting the need for a robust security framework that balances security enhancements with minimal performance impact.

According to my knowledge as of April 2024, no major work has been done at the intersection of Zero Trust and Kubernetes.

1.3 Objectives and Methodology

As there is no research done on ZT specifically in the context of Kubernetes, the objective of this thesis is to fill exactly this gap. How are ZT principles applied in Kubernetes? How can ZT be achieved in a Kubernetes setup? The aim is to construct a Kubernetes environment according to ZT principles as outlined in NIST Special Publication 800-207. To narrow down the scope of the research, the following limitations are placed.

- 1) A self-managed cluster (a self-managed control plane) will be instantiated using kubeadm. It will not be a managed cluster like EKS, AKS, or GKE. The cluster will have one master node and two worker nodes.
- 2) In order to implement ZT guidelines, open source, free, or free versions of otherwise paid-services are used. Least possible vendor lock-in is preferred.

Formulating above mentioned goals in a single, coherent sentence results in the following research question: How can free software projects and native Kubernetes solutions be leveraged to achieve a cloud-agnostic and zero-trust architecture in Kubernetes?

To do so, NIST Special Publication 800-207 is used as a primary point of reference. The paper will be thoroughly read and summarized in chapter two. Here, an in depth explanation for the importance of ZT shall be given. Core Kubernetes concepts which help understand the subsequent practical setup are also explained in chapter 2. After explaining the pillars of ZT and Kubernetes, technologies and solutions are researched and picked, which will then be practically implemented to achieve the corresponding ZT principle inside the Kubernetes environment. Code snippets shall be presented. Throughout the practical setup, cross references to the NIST publication shall be made. To evaluate the setup, test namespaces will be created and simple test applications will be deployed inside them. User accounts will be created. At the end of the paper we shall critique the environment and answer the question if a setup, which does a 100% justice to ZT, is even possible.

2 Concepts

2.1 Zero Trust

2.1.1 Definition and principles

ZT is not one single architecture, but rather a set of guidelines for more secure system architecture and operations. NIST publication presents the following definition for ZT and ZTA architecture:

"Zero trust (ZT) provides a collection of concepts and ideas designed to minimize uncertainty in enforcing accurate, least privilege per-request access decisions in information systems and services in the face of a network viewed as compromised. Zero trust architecture (ZTA) is an enterprise's cybersecurity plan that utilizes zero trust concepts and encompasses component relationships, workflow planning, and access policies. Therefore, a zero trust enterprise is the network infrastructure (physical and virtual) and operational policies that are in place for an enterprise as a product of a zero trust architecture plan"

The following principles are established. It must be noted that not every single one of them may be fully realized to a perfect degree, depending on a given use case or strategy.

- Every resource is viewed as a possible target and has to be verified as trustworthy. This includes personal devices that access company resources.
- No matter where one is located on the network, communication is encrypted and authenticated, thus trust isn't given out by default based only on network ownership.
- Access to resources is allowed on a per-session basis. Privileges are only allocated when absolutely necessary for the task at hand; access to additional resources is not granted automatically.
- Policies, created based on risk tolerance, dynamically decide access to resources by taking into account a variety of parameters, like client identity and specific application.
- The integrity and security posture of the assets are constantly monitored, actions are taken based on their security state.
- Identity, Credential, and Access Management systems provide ongoing monitoring and potential reevaluation throughout user transactions.
- Extensive data collection about asset state, network traffic, and access requests is done to have a clear picture of the current security posture of the infrastructure.

[cite NIST here] [cite british gv]

2.1.2 The network

In brief, a Zero Trust Architecture (ZTA) approach to network planning and deployment operates on these core assumptions, which closely follow above mentioned tenets.

No implicit trust

The entire enterprise network is treated as untrusted. All assets should behave as though attackers are present, necessitating secure communication through authentication and encryption.

External Device Consideration

Networked devices may not be owned or controlled by the enterprise, including those used by visitors or under BYOD policies. Authentication and authorization are required for these devices to access enterprise resources.

Trust Evaluation

The notion of inherent trust for any resource is rejected. Each asset's security posture before granting access to enterprise resources is continuously assessed, emphasizing the need for device authentication alongside subject credentials.

Infrastructure Diversity

Not all enterprise resources reside on enterprise-owned infrastructure. This encompasses cloud services and remote enterprise subjects, which may rely on external networks for connectivity.

Local Network Skepticism

A lack of trust in local network connections for remote subjects and assets is assumed. All connection requests should be treated with suspicion, requiring robust authentication and authorization measures for secure communication.

Consistent Security Policy

Consistent security policies and postures for assets and workflows transitioning between enterprise and external infrastructure are enforced. This applies to devices moving across networks and workloads migrating between on-premises and cloud environments.

[cite NIST here] [cite british gv]

2.1.3 Logical components

A ZTA deployment in an organisation consists of many logical parts. These parts can be used with a cloud-based service or as an on-premises solution. The policy engine and policy administrator (described below) are the two logical parts of the policy decision point (PDP). Application data is transferred on a data plane, whereas the ZTA logical components communicate via a separate control plane.

Policy Engine

The PE is in charge of deciding whether or not to allow access to a resource for a particular resource. To give, deny, or revoke access to a resource, the PE employs enterprise policy together with input from external sources as input to a trust algorithm. The component of the policy administrator is matched with the PE. The decision is made by the policy engine, which also documents its approval or rejection. The policy administrator then puts the decision into action.



Figure 2.1: Logical components of ZTA[source: NIST]

Policy Administrator

The PA is the component handling initiation and/or termination of the communication line between a resource and a subject. Any session-specific authentication, credential, or token that a client uses to get access to an enterprise resource would be generated by it. It is directly related to the PE and depends on its final determination of whether to approve or reject a session. The PA sets up the PEP to permit the session to begin if the request has been authenticated and the session is allowed. The PA notifies the PEP to terminate the connection in the event that the session is rejected or an earlier approval is overturned.

Policy Enforcement Point

The PEP is in charge of permitting and overseeing connections between a subject and an enterprise resource. The PEP interacts with the PA in order to transmit requests and/or obtain updates on PA policies. Although there is just one logical component in ZTA, it can be divided into two parts: the resource side (such as the gateway component in front of the resource that regulates access) and the client (such as an agent on a laptop) or a single portal component that serves as a gatekeeper for communication channels. The trust zone containing the enterprise resource is located beyond the PEP, as seen in figure 2.1.

[cite NIST here]

2.2 Kubernetes

2.2.1 Components

A Kubernetes cluster is separated into a control plane and data plane. The control plane groups the components responsible for the underlying system of Kubernetes itself. The data plane hosts the actual user containers. Machines that run control-plane services are referred to as master-nodes, whereas machines that host user applications are referred to as worker-nodes. There are four core control-plane components which work together to manage the cluster and ensure that applications run smoothly. These components run as containers themselves inside the reserved 'kube-system' namespace.

The API server in Kubernetes acts as the gateway, offering a REST API through which users and applications can communicate with the cluster, enabling resource management, access control, and querying of cluster state and configurations. The Kubelet, an agent



Figure 2.2: Kubernetes Architecture [source: author]

process on each node, ensures that pods are running healthy and up-to-date with the desired state. It does so by communicating with the API server. Etcd, a distributed key-value store, maintains the cluster's consistent state across multiple nodes, utilized by control plane components for storing cluster information like resource configurations and pod statuses. The kube-scheduler allocates pods to nodes based on resource availability and requirements, aiming to optimize successful pod execution while enhancing cluster resilience by distributing pods across multiple nodes. The controller manager, operating on each node, oversees the cluster's state, responds to events like pod creation or deletion, and does corrective measures to maintain cluster health, such as restarting unhealthy pods or scaling to meet demand. [cite k8s docs]

2.2.2 Core Cluster Objects

Namespaces

Within a single cluster, resource groups are arranged and separated using namespaces, which are virtual partitions in Kubernetes. Namespaces allow the cluster to be logically divided into smaller, more manageable chunks. Namespaces are used by Kubernetes objects like "Deployment," "Service," and "ConfigMap" to describe an application's state. Globally existing and not namespaced objects are those that define the cluster as a whole or that apply to all applications across all namespaces. Namespace isolation preserves the integrity of every deployment by averting unintentional interactions or interference between applications. Custom authorization rules and access control techniques can be enforced by assigning distinct security policies to namespaces. This aids in resource protection and limits unauthorised access between specific namespaces inside the cluster. [cite k8s docs]

Pods, Replicasets, Deployments, Statefulsets

In Kubernetes, pods are fundamental building blocks for deploying and managing containerized applications. A pod is a unit of one or more containers, which belong together and share

resources such as network namespaces and filesystems. This tight coupling allows pods to be treated as a single unit for scheduling, management, and resource allocation. Pods expose the ports exposed by its containers, allowing external traffic to reach its intended recipient. Pods are rarely instantiated by itself. Deployments, Replicasets and Statefulsets are higher deployment configurations which specify pod settings. [cite k8s docs]

ServiceAccount and RBAC

Pods can use a ServiceAccount as a security identity to access Kubernetes services and resources. It controls pod access permissions and interactions with the cluster and serves as an authentication method for pods. ServiceAccounts have the ability to mount Secrets into Pods, giving them access to private information like certificates, API keys, and passwords. RBAC is an authorization system that regulates which individuals or groups are able to access resources and carry out particular tasks inside the cluster. It offers a detailed method for controlling access rights and enforcing security regulations throughout the Kubernetes environment. Role objects define a set of permissions that can be granted to users or groups. They specify which resources can be accessed, what actions can be performed on those resources, and the scope of the permissions (namespace-wide or cluster-wide). ClusterRole objects are similar to Roles, but they have a broader scope, applying to all namespaces within the cluster. They are typically used to define permissions for system components or users who need access to resources across multiple namespaces. RoleBindings and ClusterRoleBindings assign the designated permissions to users or groups based on the association between Roles or ClusterRoles. They specify how identities and the permissions they possess are mapped out. Such RoleBindings are bound to previously mentioned ServiceAccounts. [cite k8s docs]

2.2.3 Kubernetes Networking

A key component of Kubernetes that guarantees smooth communication between containers, pods, and services inside a Kubernetes cluster is Kubernetes networking. Networking is made uniform by the Kubernetes network model.

Pod Networking

Each pod in Kubernetes has its own IP address, containers within a pod can freely communicate with one another via localhost. With these pod IP addresses, pods can communicate with any other pod in the cluster without requiring Network Address Translation (NAT). Network policies are used to establish pod isolation; pods are treated in the same way as virtual machines (VMs) or hosts with distinct IP addresses. Because they operate in the same network namespace, containers inside of pods share the IP address of the pod.

Service and Ingress

Access to a collection of pods is abstracted by Kubernetes services and is often represented by a virtual IP address within the cluster. Node ports or load balancers allow access to these services from outside the cluster. Even when the underlying pods change, a service's virtual IP address and DNS name stay the same. Load balancing and flexible access control are thus made possible.

Container Network Interface (Plugins)

The Container Network Interface (CNI) API provided by Kubernetes allows for the support of several network solutions. Third-party network implementations are increasingly utilized, even though Kubernetes' integrated network support, kubenet, provides basic connectivity. Pods can be connected to the network using network plugins, and pod IP addresses can be assigned using IP Address Management (IPAM) plugins. Calico and Cilium, for instance, offer flexibility in selecting the optimal networking solutions for certain needs and settings by providing both network and IPAM plugins and integrating with other CNI plugins.

[cite tigera, cilium, k8s docs]

3 Creating the Zero Trust Cluster

3.1 Virtual Machines and Network

The base infrastructure of the thesis project consists of 5 Virtual Machines which are running on a Proxmox infrastructure provided by FH Campus Wien. At the edge of the network sits an OPNsense gateway, an open-source firewall based on FreeBSD, which controls external traffic coming into this thesis' network. Its external interface on 10.140.0.50 is part of a broader WAN provided by institution's facilities. This gateway serves as the primary entry and exit point for all traffic entering or leaving the internal network, playing a critical role in securing and directing network traffic. The gateway runs a Wireguard agent - the general environment is accessed using SSH through a Wireguard VPN tunnel. As this thesis focuses on the Kubernetes environment, and as OPNsense is not directly part of the Kubernetes setup, there shall be no further in-depth description of this OPNsense gateway.

The LAN is segmented into a distinct subnet 192.168.20.0/24. Within this subnet, four VMs have been deployed, each assigned static IP addresses ranging from 192.168.20.1 to 192.168.20.5. These VMs serve various roles within the network, hosting different services and agents required for the Kubernetes environment and its Zero Trust character. A static IP address assignment was chosen to ensure consistent network addressing, and to facilitate easier management and configuration of these VMs. The Ubuntu installation wizard was used to complete base configurations such as username, password, hostname and IP addresses. Each of these VMs runs Ubuntu 22.04 LTS, a popular Linux distribution known for its stability, security features, and extensive software repositories. This choice of operating system emphasizes a preference for a robust, community-supported platform capable of handling the demands of modern networked applications. With 8GB of RAM and 4 CPU cores allocated to each VM, they are well-equipped to run Kubernetes.

To give a general introduction of the setup, an authentication proxy is deployed on a standalone VM on 192.168.20.2, apart from the Kubernetes cluster itself. The cluster is deployed on the remaining three VMs. 192.168.20.3, the master node, is running the control-plane (the "brain" of the cluster), consisting of kube-apiserver, etcd, kube-controller-manager and kube-scheduler. 192.168.20.4 and 192.168.20.5 are dedicated worker nodes of the cluster running the data-plane (the "body" of the cluster), hosting deployed containers. Each node runs the Kubelet and the kube-proxy. The cluster is organized in multiple namespaces. The components and software solutions which ensure a Zero Trust approach are deployed in their respective, dedicated namespaces, such as istio-system, keycloak and teleport-agent, for example. To mimic actual business software hosted on Kubernetes and for the sake of demonstration, namespace "neptune" and namespace "jupiter" are created. Each of them run simple deployments such as a Nginx webserver. This setup, as shown in figure 3.1, will be explained in depth in the following sections.



Figure 3.1: General Setup [source: author]

3.2 Bootstrapping the cluster

Before installing Kubernetes binaries on the nodes, certain settings and kernel parameters need to be configured on the Linux machines that make up the Kubernetes nodes. The bash script as seen in figure 3.2 is executed on all three nodes.

```
1  #!/bin/bash
2
3  cat <<EOF | tee /etc/modules-load.d/k8s.conf
4  overlay
5  br_netfilter
6  EOF
7
8  modprobe overlay
9  modprobe br_netfilter
10
11 # sysctl params required by setup, params persist across reboots
12 cat <<EOF | tee /etc/sysctl.d/k8s.conf
13 net.bridge.bridge-nf-call-iptables = 1
14 net.bridge.bridge-nf-call-ip6tables = 1
15 net.ipv4.ip_forward = 1
16 EOF
17
18 # Apply sysctl params without reboot
19 sudo sysctl --system
20
21 #disable swap for kubeadm to work properly
22 sudo swapoff -a
23 (crontab -l 2>/dev/null; echo "@reboot /sbin/swapoff -a") | crontab - || true
24
```

Figure 3.2: Base k8s-node configuration [source: author]

This script places a k8s.conf file under the /etc/modules-load.d/ directory, specifying overlay and bridge netfilter. This ensures that the kernel modules overlay and bridge netfilter are always loaded at startup of the VM, which are required for Kubernetes networking and

storage requirement.

Overlay

The overlay module is essential for Kubernetes because it activates the overlay storage driver used by container runtimes such as Docker to manage container filesystems efficiently. This driver allows containers to share and layer filesystems, optimizing storage use and enhancing performance. By enabling the overlay module, Kubernetes ensures that container runtimes can efficiently create, manage, and run containers, which is crucial for the effective functioning of the cluster.

Bridge Netfilter

The bridge netfilter module is important because it allows the Linux bridge to interface with the iptables firewall, enabling packet filtering and NAT. This capability is needed for Kubernetes to manage and route network traffic between pods, services, and external networks efficiently. `net.bridge.bridge-nf-call-iptables` and `net.bridge.bridge-nf-call-ip6tables` settings ensure that bridged traffic goes through the iptables chains. `net.ipv4.ip-forward` enables IP forwarding, allowing the Linux kernel to forward packets between (virtual) network interfaces.

With the `modprobe` command, these modules are loaded into the kernel and their functions are made available. Without these settings, Kubernetes components may experience problems with container filesystem management, network packet processing, and general traffic routing.

Swap

Disabling swap on Linux during Kubernetes setup, as done in lines 22-23 of figure 3.2, is necessary because Kubernetes depends on efficient memory management for container orchestration. When swap is active, the kernel might swap container memory to disk. This results in unpredictable performance and potential container instability. By turning off swap, Kubernetes ensures that container memory stays in physical RAM. This allows Kubernetes to manage containers effectively without interference from swapping mechanisms.

Installation

After aforementioned settings, one can proceed with installing the container runtime, `kubeadm`, `kubelet`, and `kubectl` on all nodes. `kubeadm` is then used to initialize the cluster. `kubectl` is used to interact with the control plane of the cluster. The script as in figure 3.3 downloads GPG keys, installs and runs CRI-O. Figure 3.4 shows installation of Kubernetes binaries. Finally the control plane is initialized on the master node (192.168.20.3), as seen in figure 3.5, and worker nodes (192.168.20.4 and 192.168.20.5) are joined to the cluster, as seen in figure 3.7.

3 Creating the Zero Trust Cluster

```
1  #!/bin/bash
2
3  #download GPG keys and install CRI-O container runtime
4  apt-get update -y
5  apt-get install -y software-properties-common curl apt-transport-https ca-certificates
6
7  curl -fsSL https://pkgs.k8s.io/addons:/cri-o:/prerelease:/main/deb/Release.key |
8  gpg --dearmor -o /etc/apt/keyrings/cri-o-apt-keyring.gpg
9  echo "deb [signed-by=/etc/apt/keyrings/cri-o-apt-keyring.gpg] https://pkgs.k8s.io/addons:/cri-o:/prerelease:/main/deb/ /*
10  tee /etc/apt/sources.list.d/cri-o.list
11
12  apt-get update -y
13  apt-get install -y cri-o
14
15  systemctl daemon-reload
16  systemctl enable crio --now
17  systemctl start crio.service
```

Figure 3.3: Container runtime installation [source: author]

```
1  #!/bin/bash
2
3  #download GPG keys and install components
4  KUBERNETES_VERSION=1.29
5  mkdir -p /etc/apt/keyrings
6  curl -fsSL https://pkgs.k8s.io/core:/stable:/v$KUBERNETES_VERSION/deb/Release.key | \
7  gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
8  echo "deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v$KUBERNETES_VERSION/deb/ /* | \
9  tee /etc/apt/sources.list.d/kubernetes.list
10  apt-get update -y
11  apt-get install -y kubelet kubeadm kubectl
```

Figure 3.4: Kubernetes components installation [source: author]

```
1  #!/bin/bash
2  kubeadm init --pod-network-cidr=10.0.0.0/24 --apiserver-advertise-address 192.168.20.3
```

Figure 3.5: Control plane initialization [source: author]

```
1  #!/bin/bash
2  kubeadm join 192.168.20.3:6443 --token [token from kubeadm init] --discovery-token-ca-cert-hash [token from kubeadm init]
```

Figure 3.6: Join nodes to cluster [source: author]

Having securely downloaded and installed dependencies, the following sections discuss fortifying the cluster according to Zero Trust principles.

3.3 Teleport

3.3.1 Why use Teleport?

Kubernetes clusters are administered and accessed by various different professionals: security engineers, database administrators, software developers, DevOps engineers, general system administrators and others. Commandline clients such as 'kubectl' or 'helm' are used to manage the cluster. Such access must happen in a secure, granular and audited fashion in order for it to fulfill ZT requirements. The default way of accessing the cluster is done by using static client keys and certificates placed under `.kube/config` of the client machine. However, access to the cluster can also happen by means of a dedicated authentication-proxy solution. cite[k8s]

Kubeconfig, Teleport, Okta and Hoop are ways of accessing the cluster which were researched for suitability for the setup of this thesis. Kubeconfig means accessing the cluster through a manually issued kubeconfig file. Teleport, Okta and Hoop were identified as authentication-proxy solutions to the cluster, aiming to implement ZT. To identify the most appropriate solution for a ZT approach, following criteria was established:

- Kubernetes friendly
- Is free and open source software (FOSS)
- Centralized identity management taking care of rotating secrets
- Offers Role Based Access Control (RBAC)
- Offers Attribute Based Access Control (ABAC)
- Provides MFA
- Possibility to audit sessions

Out of the box, Kubernetes offers its standard kubeconfig file, consumed by kubectl and helm, to manage Kubernetes clusters. This is inherently Kubernetes friendly, free and open source. This method offers RBAC as part of Kubernetes' integrated RBAC implementation, achieved through Role, ClusterRole, RoleBinding and ClusterRoleBinding resources. This method fundamentally lacks the features needed to come closer to a ZT approach. Kubernetes itself does not offer one centralized way of creating, issuing and rotating these credentials. They are permanently valid until they are manually invalidated, creating a long-lived, single point of compromise. For this reason, authentication proxies such as Teleport, Okta, and Hoop are more appropriate solutions for cluster access. Okta is a mature identity and access management platform offering robust security features for ZT. It can act as an access gateway in front of Kubernetes clusters. However, this feature is not free and open source software and cannot be self-hosted. It was therefore not taken into further consideration. Hoop.dev and Teleport remain as promising solutions. Apart from paid subscriptions they both offer free, self hosted versions and therefore qualify for the setup. The decision for Teleport comes down to the fact that Teleport natively provides even more granular access control through ABAC, implemented by Teleport predicate language. The comparison is summarized in figure 3.7.

3 Creating the Zero Trust Cluster

	k8s-friendly	FOSS	Centralized	RBAC	ABAC	MFA	Auditing
Kubeconfig	yes	yes	no	yes	no	no	no
Teleport	yes	yes	yes	yes	yes	yes	yes
Okta	no	no	yes	yes	yes	yes	yes
Hoop	no	yes	yes	yes	no	yes	yes

Figure 3.7: Comparison-summary of access solutions to cluster [source: author]

With Teleport, critical features are implemented, which help shaping the cluster towards Zero Trust. This includes Identity-Driven Authentication, elimination of static credentials, RBAC, ABAC and MFA.

Identity-Driven Authentication and Authorization

Teleport acts as a gateway for accessing corporate resources, both on-premises and in the cloud, by leveraging identity-based authentication and authorization. This method eliminates the need for traditional VPNs and password management tools, focusing instead on verifying the identity of users rather than their network location or credentials. This approach is central to the Zero Trust philosophy, which grants access based on the identity and purpose of the user, not their physical location or network status, as aforementioned in section 2.1. cite[teleport]

Elimination of Static Credentials

Teleport issues short-lived identity certificates instead of static credentials. This practice minimizes the risk associated with credential theft, as these certificates automatically expire after a certain period, contrasting sharply with the long-term validity of static credentials. This strategy is consistent with the Zero Trust ethos of minimizing vulnerabilities by avoiding the use of persistent secrets. cite[teleport]

Role and Attribute Based Access Control and Auditing

Teleport implements role-based access control, allowing access to resources within the corporate infrastructure based on the roles assigned to individual users' identities. Additionally, it can record all sessions and actions, This provides organizations with the ability to monitor and trace access to their infrastructure securely. Such level of control and oversight is crucial for implementing ZT principles, which advocate for continuous verification of access rights and thorough monitoring for potential security threats, as mentioned in NIST800-207, and as summarized in chapter 2. cite[teleport]

Public-Facing Proxy and Just-In-Time Access

Teleport runs a public-facing proxy service that facilitates secure access to resources without the necessity of a VPN. This feature enables just-in-time access, where users are granted temporary privileges to access specific resources or roles as required. This method strengthens security by ensuring that users and devices are only granted access based on their immediate needs, reducing the likelihood of unauthorized access or lateral movement within the network. cite[teleport]

Multi-Factor Authentication

Multi-Factor Authentication (MFA) is a security measure that requires users to provide two or more verification factors to access a resource like an application or online account. These factors can include something the user knows (like a password), something the user has (such as a security token or a smartphone), or something the user is (fingerprint, facial recognition). By requiring multiple types of verification, MFA significantly increases the difficulty for unauthorized individuals to gain access, even if one factor is compromised. This method is a crucial part of Identity and Access Management policies, aiming to enhance security against cyber threats by adding layers of authentication beyond just a username and password. cite[aws]

3.3.2 Installation and setup

Teleport runs on a separate VM on 192.168.20.2, next to the Kubernetes nodes. The script shown in figure 3.7...

1. Fetches teleports public PGP key
2. Adds Teleport repository to APT package manager
3. Installs Teleport
4. Generates certificate and private key for HTTPS GUI access
5. Configures and runs Teleport

```
1  #!/bin/bash
2
3  #Download Teleport's PGP public key
4  curl https://deb.releases.teleport.dev/teleport-pubkey.asc \
5      -o /usr/share/keyrings/teleport-archive-keyring.asc
6
7  #Add the Teleport APT repository
8  echo "deb [signed-by=/usr/share/keyrings/teleport-archive-keyring.asc] https://deb.releases.teleport.dev/ stable main" \
9      | tee /etc/apt/sources.list.d/teleport.list > /dev/null
10
11 apt-get update
12 apt-get install teleport
13
14 #generate cert for https GUI access
15 openssl genpkey -algorithm RSA -out teleportkey.pem
16 openssl req -new -key teleportkey.pem -out teleportcsr.pem
17 openssl x509 -req -days 365 -in teleportcsr.pem -signkey teleportkey.pem -out teleportcert.pem
18 cp teleportkey.pem /var/lib/teleport/privkey.pem
19 cp teleportcert.pem /var/lib/teleport/fullchain.pem
20
21 #configure teleport
22 teleport configure --acme --acme-email=guntram.klaus@stud.fh-campuswien.ac.at --cluster-name=teleport.nebhis.lcl \
23     --cert-file=/var/lib/teleport/fullchain.pem --key-file=/var/lib/teleport/privkey.pem | tee /etc/teleport.yaml
24
25 #run teleport
26 systemctl enable teleport
27 systemctl start teleport
```

Figure 3.8: Teleport installation [source: author]

Authy

Authy is a mobile application and service that offers two-factor authentication (2FA) for online accounts. Authy is designed to be compatible with multiple devices, enabling users to synchronize their 2FA tokens across different devices for added convenience and security. Teleport requires integration to MFA from the very start. When accessing the Teleport UI as well as when using the tsh commandline client, users are prompted for One-Time password tokens (OTP).

Architecture

Figure 3.9 provides an overview of the Teleport setup. The key components are the user, the Teleport proxy, the Teleport agent, and finally the Kubernetes API server. The user has "tsh" installed on their local machine, which is a commandline client used to interact with a given Teleport server. By running "tsh login -proxy=teleport.nebhis.lcl -user=teleport-admin", the user "teleport-admin" authenticates to the Teleport cluster "teleport.nebhis.lcl", hosted on 192.168.20.2. This Teleport proxy acts as an intermediary in front of the Kubernetes cluster, enforcing MFA and mTLS. Running "tsh kube login nebhis-kubernetes" puts the registered cluster "nebhis-kubernetes" in the current context of the kubectl command-line client by modifying the kubeconfig file under the .kube/config directory of the client machine. Subsequently, any Kubernetes API requests executed with kubectl or helm are proxied through the Teleport proxy. The Teleport Proxy establishes an encrypted tunnel to the Teleport agent running on a worker node within the "teleport-agent" namespace inside the Kubernetes cluster. Finally, the Teleport agent forwards requests to the Kubernetes API Server, which ultimately processes the requests and manages cluster resources. The API server is running on the master node inside the "kube-system" control plane namespace. This means that the Teleport proxy authenticates the actions "tsh login" and "tsh kube login", and then establishes an encrypted tunnel to the teleport-agent which issues the kubectl command to the API. The respective responsibilities of the proxy and the agent are:

Teleport Proxy:

- Authentication and Authorization, check if user is eligible for Kubernetes access
- Reverse proxying for Kubernetes API requests
- Audit logging and traffic encryption

Teleport Agent:

- Registering the Kubernetes cluster with the proxy
- Manage service accounts and their credentials for Teleport to interact with the Kubernetes API
- Synchronize configurations between the Kubernetes cluster and the Teleport Proxy

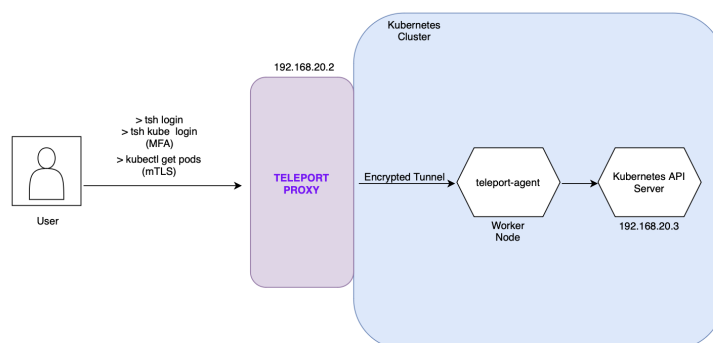


Figure 3.9: Teleport Kubernetes Setup [source: author]

Thus, access to the cluster covers most of ZT principles as described in NIST800-207.

3.4 Istio

3.4.1 Service Meshes

In a microservices architecture, a service mesh is a special infrastructure layer integrated into an application that manages service-to-service communication. It finds other services, encrypts data, manages the flow of service requests to other services, and does load balancing. It does so by injecting a so called proxy sidecar-container into Kubernetes pods. These sidecars intercept any incoming and outgoing traffic of the actual app container. That way, any logic driving matters of network communication is taken out of the app container and handled on its own plane of infrastructure. These sidecars deployed alongside app containers comprise the data-plane, while the control plane instance usually resides in its own dedicated namespace. The main benefits of a sidecar approach are sophisticated routing capabilities, separation of concerns and enhanced security. cite[redhat] cite[istiodocs] cite[dynatrace]

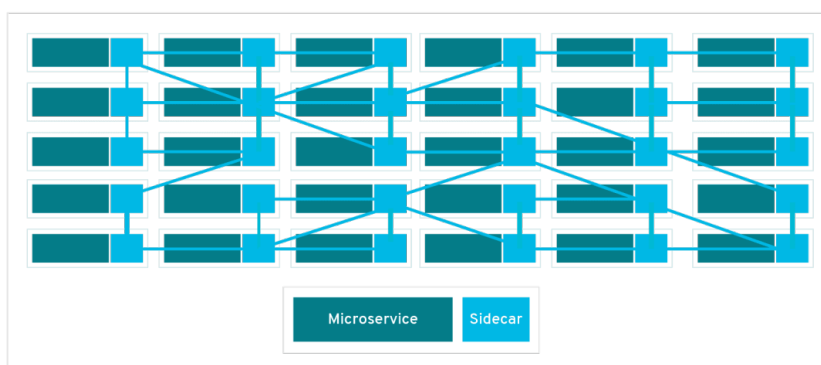


Figure 3.10: Service Mesh visualization [source: RedHat]

3.4.2 mTLS

Transport Layer Security (TLS) is a widely-used encryption protocol that authenticates servers and encrypts communication between clients and servers, ensuring data privacy and integrity. It operates using public key cryptography, which involves a pair of keys: a public key for encryption and a private key for decryption. Mutual TLS (mTLS) extends this concept by requiring both the client and server to authenticate each other using their respective TLS certificates. This mutual authentication adds a layer of security, verifying the identities of both parties and enhancing the protection against various cyber attacks. This makes it especially useful in Zero Trust architectures. Since service meshes handle network communication between pods, they must have the possibility for mTLS to be considered an employable component for a ZT Kubernetes cluster.

3.4.3 Why use Istio?

There are several ways of implementing a service mesh and there are different technologies to realize a service mesh on top of Kubernetes. Linkerd, Consul, Istio and Cilium Service Mesh are mature, free, and open source software solutions which were taken into consideration for this thesis. The decision was based on the following criteria.

- Granularity of traffic routing
- Option for mTLS

3 Creating the Zero Trust Cluster

- OIDC integration
- Audit logging
- Access logging

According to the documentation, mTLS and access logging is provided by all solutions. Audit logging is not natively supported by any of the four solutions. Therefore it comes down to the degree of granularity offered by the service meshes. To evaluate granularity, adjectives 'weak', 'limited' and 'strong' are used. Granularity of traffic routing pertains to how detailed cluster traffic can be controlled as it moves from its origin pod to its destination pod. For traffic routing, each service mesh solution provides its set of resources with which rules are defined. Controlling traffic in Linkerd is achieved through a combination of the resources ServiceProfile, TrafficSplit, Server and ServerAuthorization. Through these definitions, Linkerd allows to regulate traffic only based on service's mTLS identity. It does not offer regulating traffic based on definite IP addresses, native Kubernetes service accounts, namespaces, HTTP method or OIDC credentials. Consul leverages its service-intentions, service-resolvers, service-splitters and service-routers resources which allow for blocking traffic based on IP address and HTTP method. However, it does not natively understand underlying namespaces, service accounts and OIDC credentials. Cilium's service mesh is special as it does not inject sidecars. Rather, it leverages eBPF Linux kernel, which allows for running custom modules in kernel space of the node. This eBPF program, the Cilium agent, implements rules specified in CiliumNetworkPolicy resources. Cilium natively supports the concept of Kubernetes namespaces and service accounts. It is possible to allow or deny traffic based on namespace, IP address, service account and HTTP method. However, to validate OIDC credentials, Cilium requires additional integration and configuration with other technologies. Istio was found to support all these factors. Istio's VirtualService, DestinationRule, AuthorizationPolicy and RequestAuthentication resources enable traffic controlling also based on provided JWT tokens of the OIDC protocol, on top of other factors such as HTTP method and namespace. Linkerd and Cilium's eBPF mesh triumphs in terms of performance and efficiency, but not in terms of granularity and security. The degree of granularity offered by Istio could not be found in the other mentioned projects. It is therefore attributed with 'strong' granularity while Consul, Cilium and Linkerd are marked as 'Limited'. As such, Istio is the chosen solution for the setup. The findings are summarized in figure 3.10.

	Granularity	mTLS	OIDC	Audit Logs	Access Logs
Linkerd	Limited	yes	no	no	yes
Consul	Limited	yes	no	no	yes
Istio	Strong	yes	yes	no	yes
Cilium	Limited	yes	no	no	yes

Figure 3.11: Comparison-summary of service-mesh solutions [source: author]

3.4.4 Installation and setup

3.5 Keycloak

3.5.1 Identity and Access Management

IAM refers to all policies, processes and technologies used control access to digital resources. Through IAM systems, it is ensured that only appropriate clients have access to resources under right circumstances. For this thesis, IAM is implemented for authorized and granular pod to pod communication. In this context, the term "client" refers to a service issuing a HTTP request to another service hosted in the cluster. Using a IAM solution, the following things are taken care of.

- Identity Management: Creating client identities and their attributes.
- Authentication: Verifying the identity of clients before granting access.
- Authorization: Determining what resources an authenticated client or device can access and what actions they can perform.
- Access Control: Policies and rules enforcing which clients can access which resources under which conditions.
- Auditing: Logging client access and activities to ensure compliance and security.

The main drivers of IAM are the concepts of role-based and attribute-based access controls and the protocols Open ID Connect as well as OAuth2.

3.5.2 RBAC and ABAC

Role-Based Access Control (RBAC) is a commonly implemented access control model that allocates permissions to users based on their organizational roles. In RBAC, roles are defined by job functions, and users receive access to resources and actions that correspond to their roles. This model simplifies administration by enabling system administrators to manage permissions via roles rather than on an individual user basis. For instance, a software developer might have a role that permits access to certain repositories, whereas a database administrator might have a role that allows access to certain databases and system maintenance tools. The main benefit of RBAC is its ease of implementation and scalability in organizations with clear hierarchical structures.

Attribute-Based Access Control (ABAC) is a more adaptable and flexible access control model that grants access based on a combination of attributes pertaining to the user, the resource, and the environment. These attributes can include user details (such as department, job title, or clearance level), resource characteristics (such as data sensitivity or classification), and contextual factors (such as time of day or location). ABAC policies leverage these attributes to create intricate, fine-grained rules for access control. For instance, an ABAC policy might permit a system engineer to access certain servers only during business hours and exclusively from within the corporate network. This degree of granularity makes ABAC effective in environments where access requirements are complex and highly variable.

ABAC is more suitable for ZT because it provides the flexibility and granularity needed to implement dynamic and context-aware access controls essential for ZT. Unlike RBAC, which is relatively static and role-centric, ABAC can evaluate multiple attributes in real-time to determine access permissions, ensuring that only authenticated and authorized users can access resources under specific conditions. This aligns with ZT's core tenet of minimizing trust and

continuously verifying every access request based on a comprehensive set of criteria, thereby enhancing security and reducing the risk of unauthorized access.

3.5.3 OIDC and OAuth2

OpenID Connect (OIDC) and OAuth 2.0 are complementary protocols that together provide a framework for authentication and authorization on the web. OAuth 2.0 is the de facto standard authorization framework that allows third-party applications to obtain limited access to a web service, either on behalf of a resource owner or by allowing the third-party application to access its own resources. It works by issuing tokens that represent the user's delegated permissions, enabling secure access without sharing passwords. OpenID Connect, built on top of OAuth 2.0, extends this protocol by adding an identity layer, allowing applications to verify the identity of the end-user based on the authentication performed by an authorization server. OIDC standardizes the way identity information is exchanged using RESTful APIs, providing user profile information through an ID token. This integration of OIDC and OAuth 2.0 ensures that while OAuth 2.0 handles authorization and secure resource access, OIDC handles user authentication and identity management.

Pod to pod OIDC flow

In the context of this thesis, a service running inside the Kubernetes cluster may only access another service if the request contains a valid OIDC token, as done through the following steps.

1. The client service sends a POST request to the authorization server's token endpoint. This request contains the clientID, a client secret, and the grant type.
2. The authorization server validates the client secret and issues a short-lived access token.
3. The client receives the token and adds it to subsequent requests to the resource server. The token is added in the 'Authorization' HTTP header.
4. The resource server validates the access token by verifying its signature, using the public key obtained from the PKI endpoint of the authorization server. If the token is valid, the resource server returns the requested resource.

However, in this thesis' setup, the validation of the access token is not done by the resource service itself. This step is done by the Istio service mesh. The Istio sidecar-proxy in every application pod abstracts this handling of OIDC tokens. That way, a request does not even reach the resource server in case of invalid credentials.

3.5.4 Why use Keycloak?

Authentik, Keycloak, Gluu and Okta are four IAM solutions which were inspected for the following criteria.

- Free and open source software (FOSS)
- OAuth2 support
- OIDC support
- MFA support

3 Creating the Zero Trust Cluster

- Degree of granularity
- Option for auditing

Based on the documentations, all four solutions have comprehensive offerings for above mentioned criteria. OAuth2, OIDC, MFA and extensive granularity have become non-negotiable features of proper IAM solutions and are achieved in equal measure across Authentik, Okta, Keycloak and Gluu. Out of these, there is no product which unambiguously surpasses the other when it comes to the requirements of this setup specifically. Nonetheless, Keycloak was chosen because of its substantial community support on GitHub issues and StackOverflow, its great Kubernetes compatibility due to officially provided Helm charts and its clear, extensive documentation.

	FOSS	OAuth2	OIDC	MFA	Granularity	Auditing
Authentik	yes	yes	yes	yes	Strong	yes
Keycloak	yes	yes	yes	yes	Strong	yes
Okta	no	yes	yes	yes	Strong	yes
Gluu	yes	yes	yes	yes	Strong	yes

Figure 3.12: Comparison-summary of IAM solutions [source: author]

3.5.5 Installation and setup

RequestAuthentication, AuthorizationPolicy, strict TLS, examples with Neptune and Jupiter settings, python script: "With Istio's service mesh capabilities and Keycloak as an identity provider running in the cluster setup, tightly regulated traffic controls with aforementioned AuthorizationPolicy and RequestAuthentication policies can be achieved. EXPLAIN HOW KEYCLOAK AND ISTIO SERVE AS PDP/PEP.

3.6 Vault Secret Management

3.6.1 Why use Vault?

3.7 Trivy

3.7.1 Why Trivy?

4 Discussion and Future Work

5 Conclusion

Bibliography

List of Figures

2.1	Logical components of ZTA[source: NIST]	6
2.2	Kubernetes Architecture [source: author]	7
3.1	General Setup [source: author]	11
3.2	Base k8s-node configuration [source: author]	11
3.3	Container runtime installation [source: author]	13
3.4	Kubernetes components installation [source: author]	13
3.5	Control plane initialization [source: author]	13
3.6	Join nodes to cluster [source: author]	13
3.7	Comparison-summary of access solutions to cluster [source: author]	15
3.8	Teleport installation [source: author]	16
3.9	Teleport Kubernetes Setup [source: author]	17
3.10	Service Mesh visualization [source: RedHat]	18
3.11	Comparison-summary of service-mesh solutions [source: author]	19
3.12	Comparison-summary of IAM solutions [source: author]	22

List of Tables

Appendix

(Hier können Schaltpläne, Programme usw. eingefügt werden.)